auditor + DALEX

a powerful duet for validation and explanation of machine learning models

Alicja Gosiewska
Mi2DataLab
PhD Student at Warsaw University of Technology

Tomasz Mikołajczyk
Mi2DataLab, Ministry of Health

Instructions

1. Find the repository:

bit.ly/auditorWhyR

or

https://github.com/agosiewska/auditor-whyr2019

2. clone repository to the local folder or download zip file
3. open `part_1_explain.R` file

Do we need XAI? 1/3

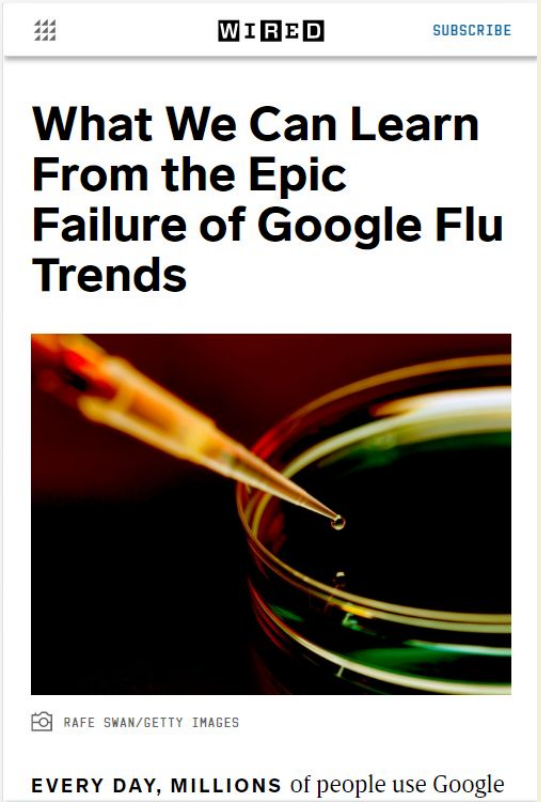

Google Flu Trends

Researchers from Google claimed that they could predict outbreaks of flu based on people's searches (paper stated that the Google Flu Trends obtains 97% accuracy).

BUT
Subsequent reports asserted that Google Flu Trends' predictions have sometimes been very inaccurate

- in the 2012-2013 flu season predicted twice as many doctors' visits as it was recorded

By re-assessing the original model, it was uncovered that the model was aggregating queries about different health conditions.

https://www.wired.com/2015/10/can-learn-epic-failure-google-flu-trends/

Do we need XAI? 2/3

IEEE SPECTRUM

Diagnostics

How IBM Watson Overpromised and Underdelivered on AI Health Care

By **ELIZA STRICKLAND**
Posted 02 Apr 2019 | 15:00 GMT

IBM Watson

Researchers from IBM took a collection of patient symptoms and came up with a list of possible diagnoses, each annotated with confidence level and links to supporting medical literature.

BUT

Watson's predictions did not earn the trust of physicians.
Watson did not achieve as good results as expected.

https://spectrum.ieee.org/biomedical/diagnostics/how-ibm-watson-overpromised-and-underdelivered-on-ai-health-care

Do we need XAI? 3/3

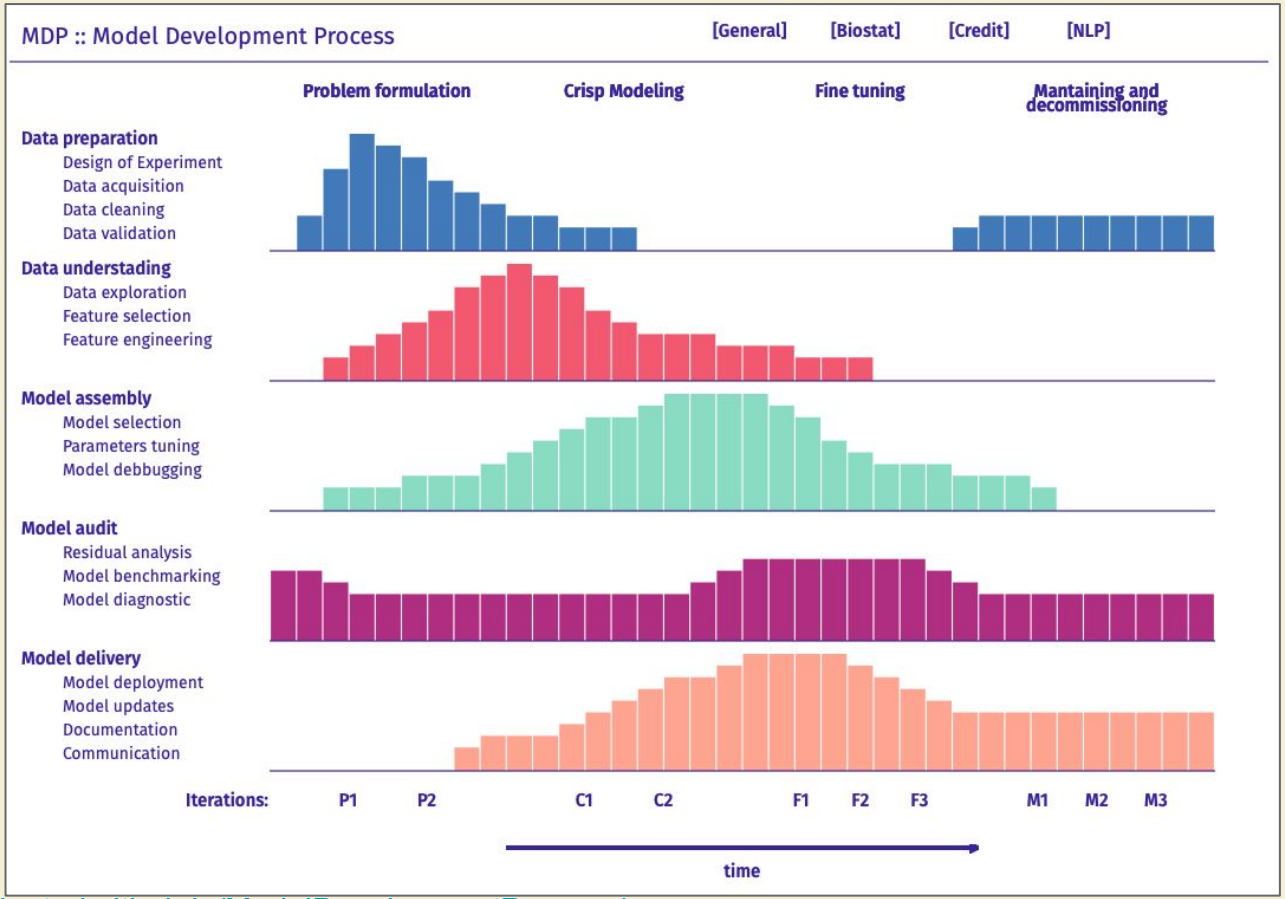

https://modeloriented.github.io/ModelDevelopmentProcess/

mlr

H2O

scikit learn

tidymodels

...

DALEX

auditor

iBreakDown

ingredients

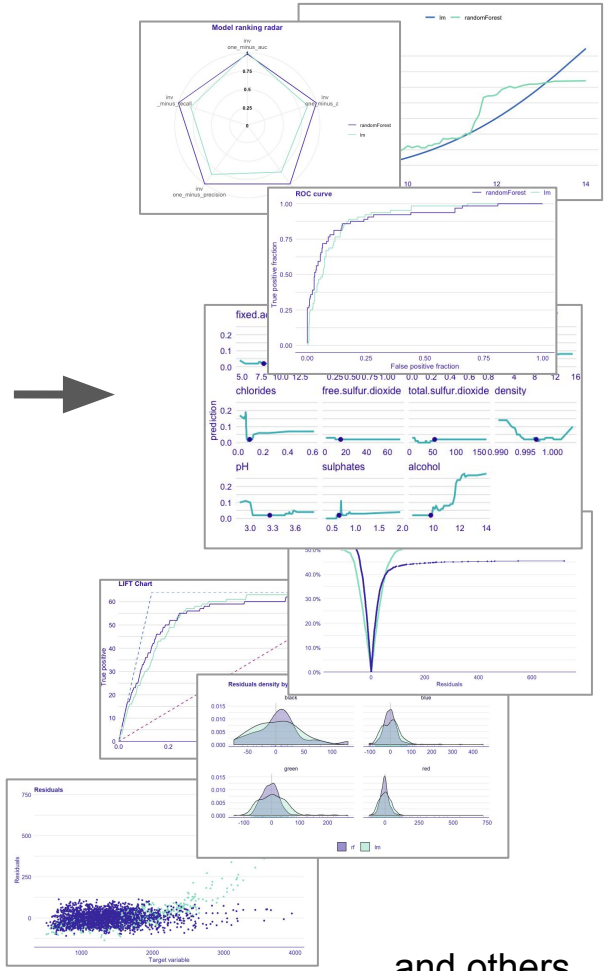

and others...

Workflow

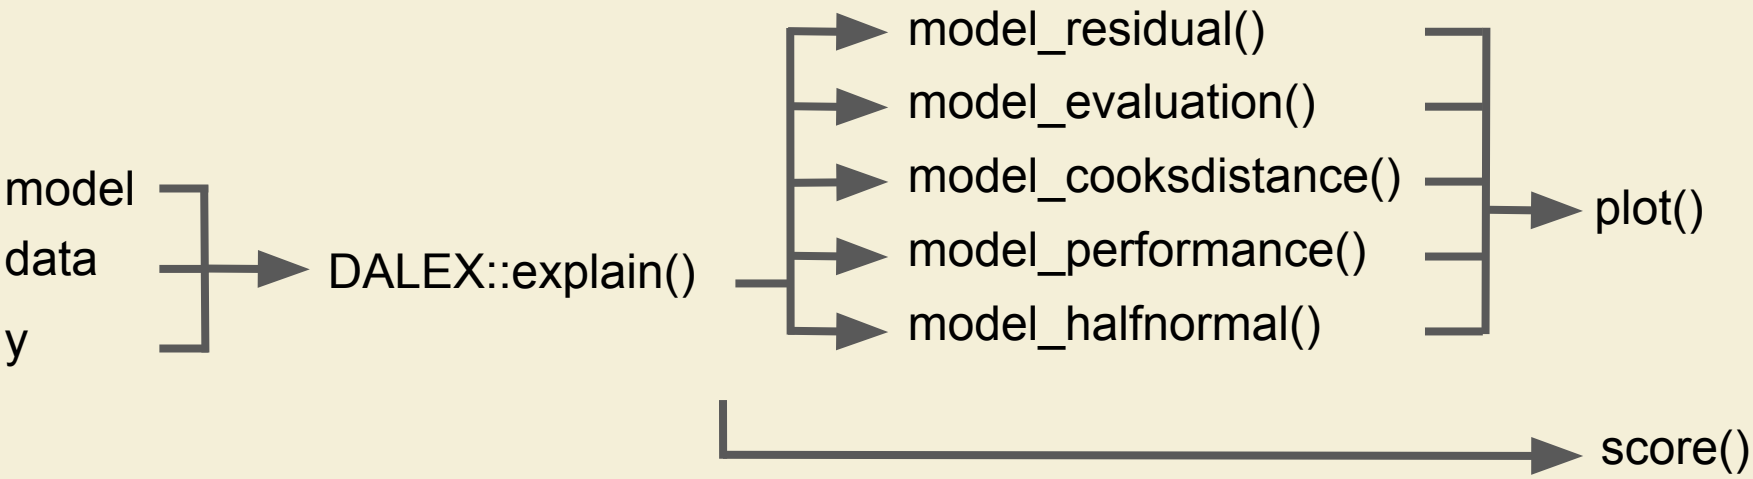
model
data
y
DALEX::explain()
model_residual()
model_evaluation()
model_cooksdistance()
model_performance()
model_halfnormal()
plot()
score()

Workflow

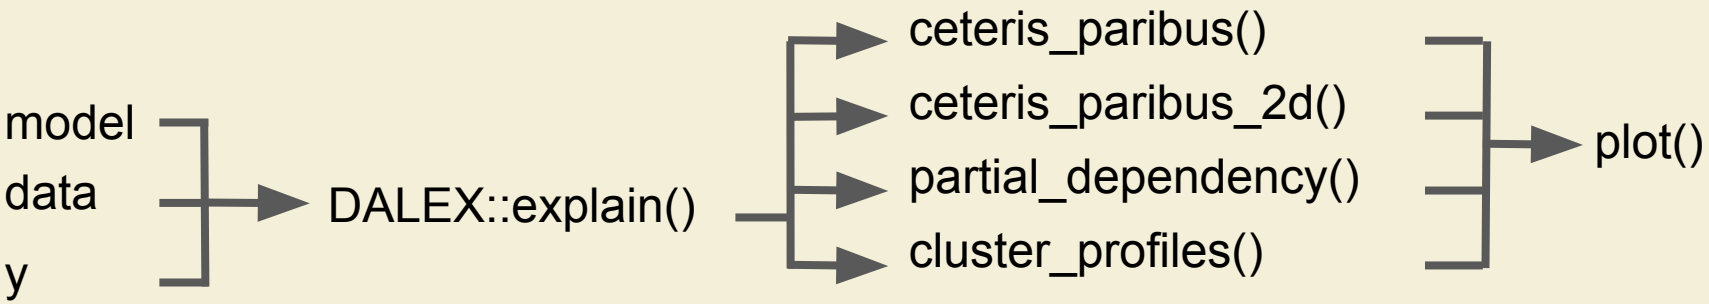
model
data
y
DALEX::explain()
ceteris_paribus()
ceteris_paribus_2d()
partial_dependency()
cluster_profiles()
plot()

Workflow

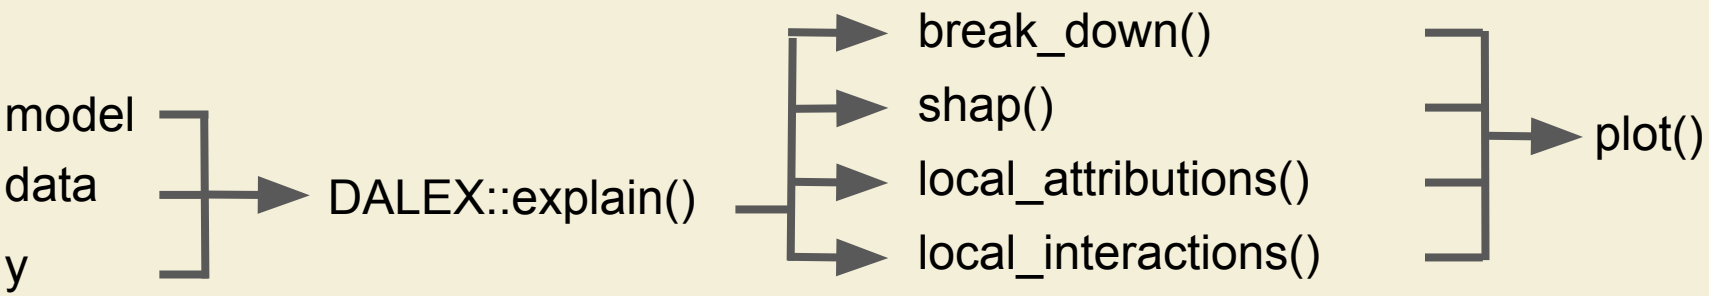
model
data
y
DALEX::explain()
break_down()
shap()
local_attributions()
local_interactions()
plot()

Plots

Radar plot

`plot_radar()`

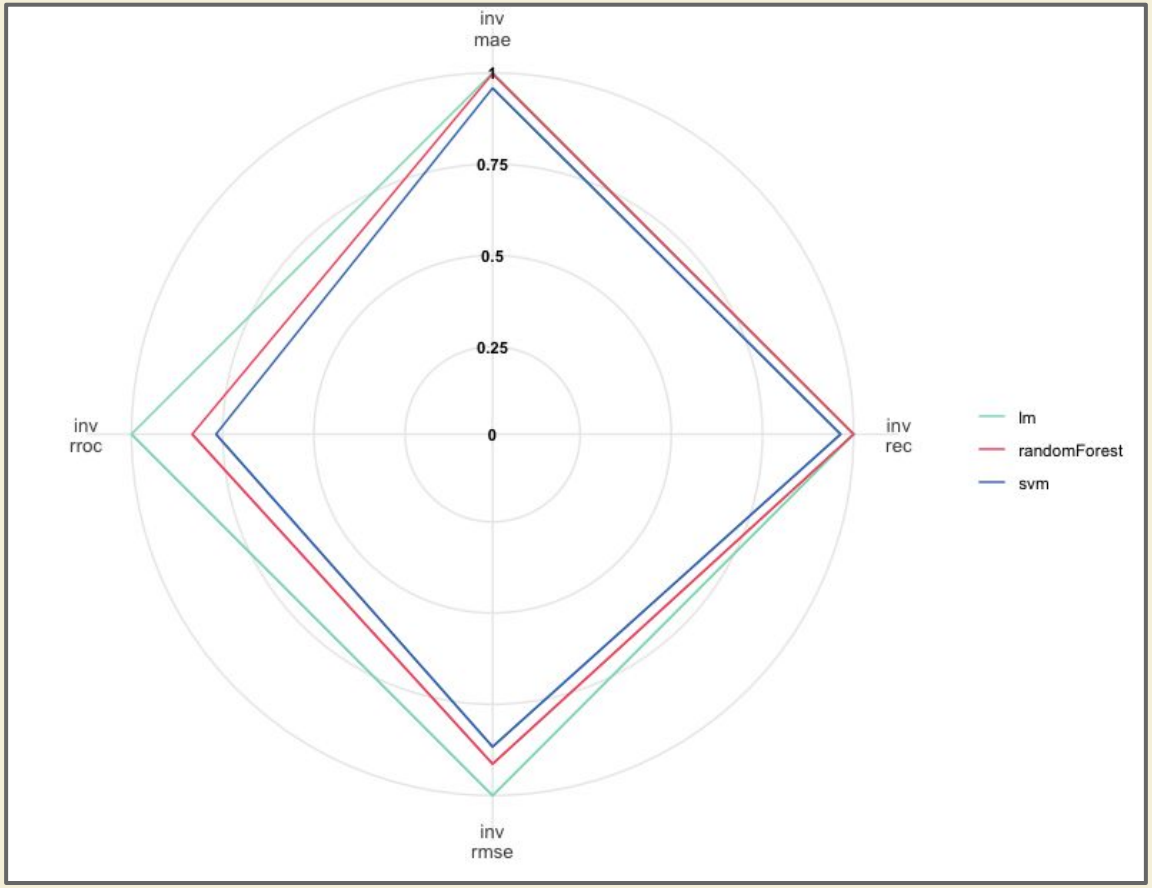

Radar plot

`plot_radar()`

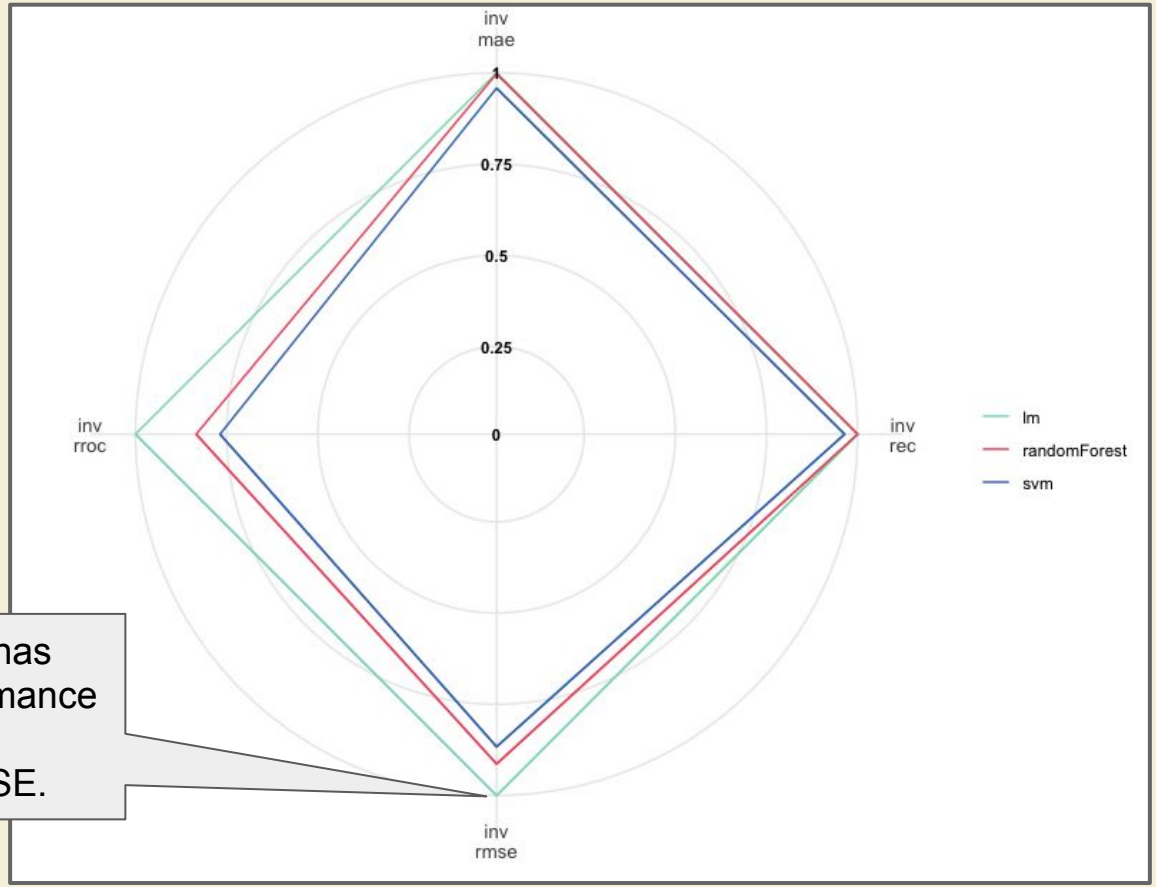

The linear model has the largest performance when models are measured by RMSE.

Receiver operating characteristic (ROC)

`plot_roc()`

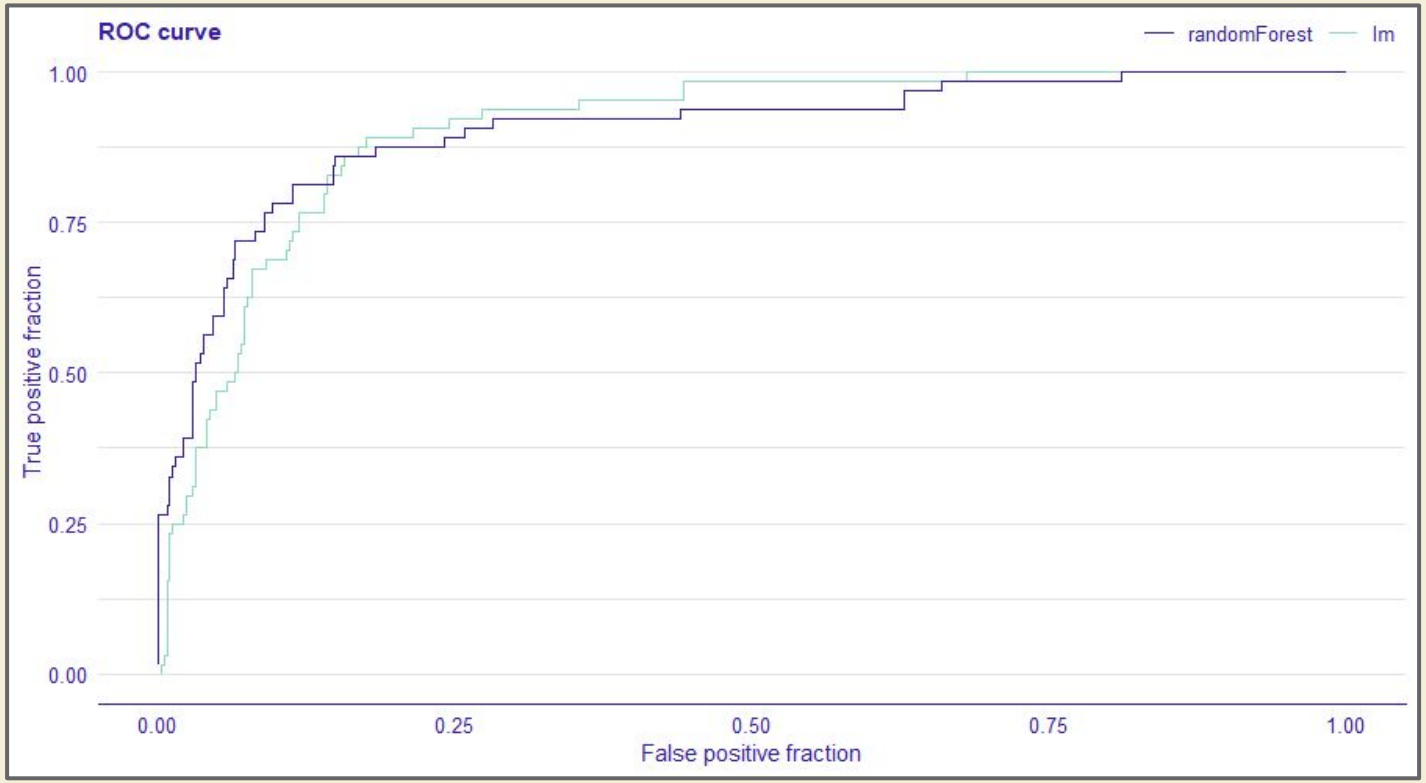

Receiver operating characteristic (ROC)

`plot_roc()`

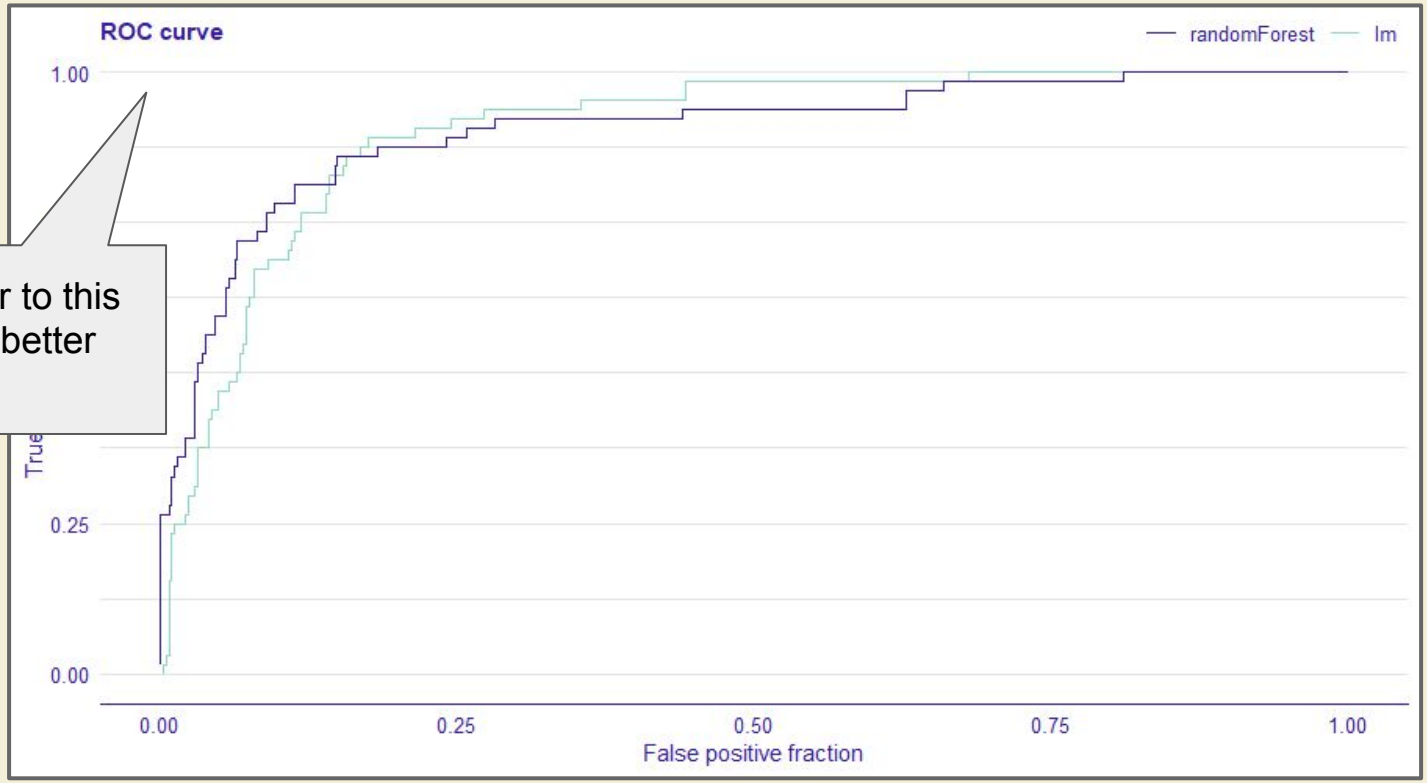

The closer to this point, the better model.

LIFT (toy example)

`plot_lift()`

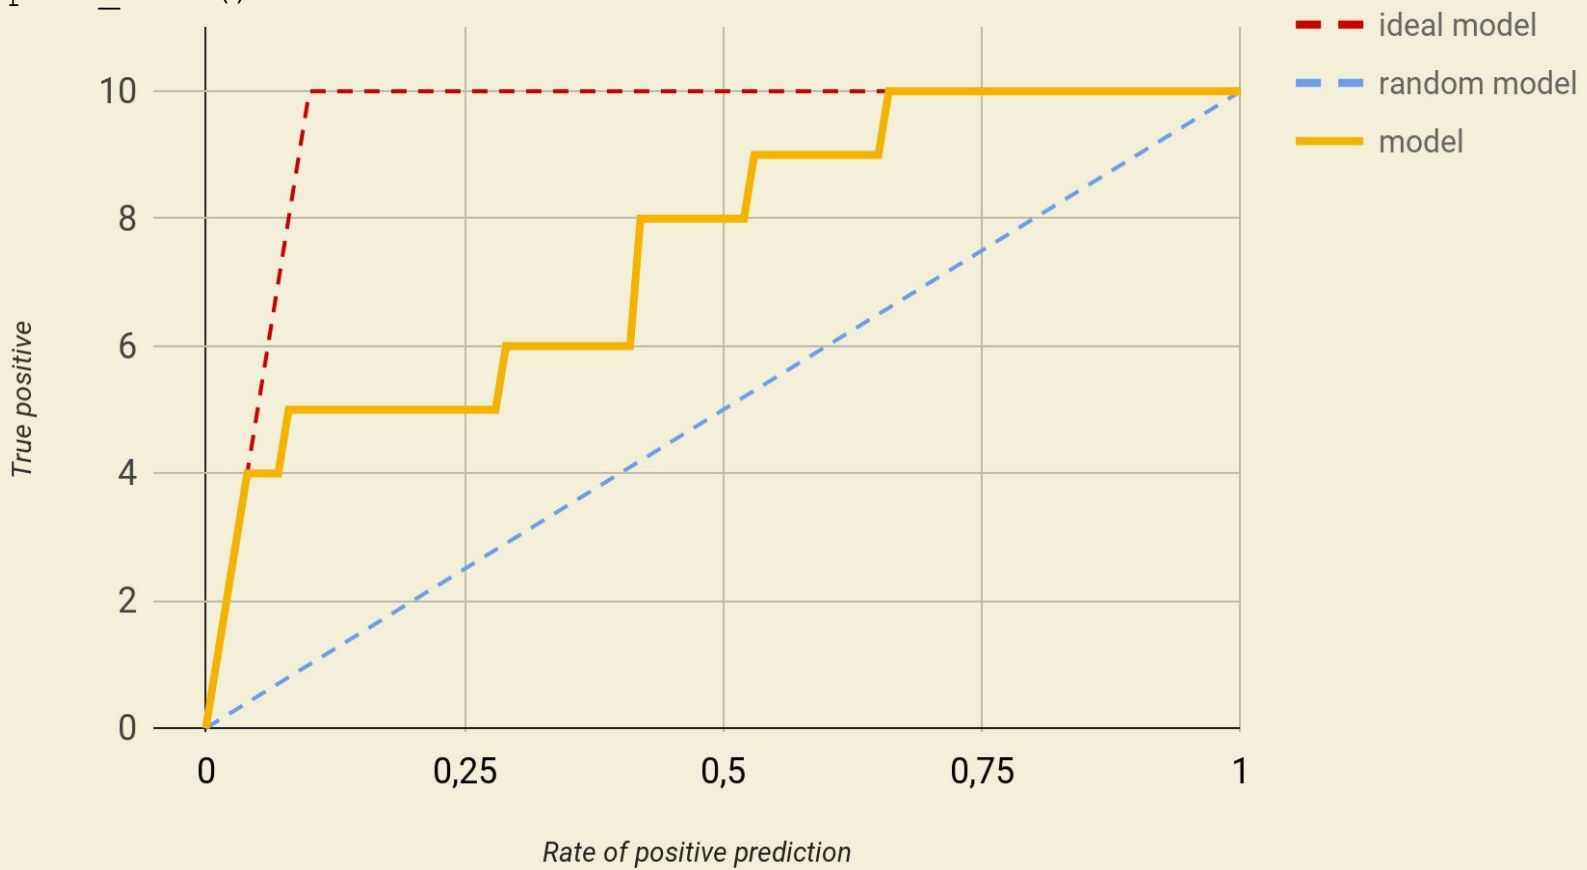

LIFT (toy example)

`plot_lift()`

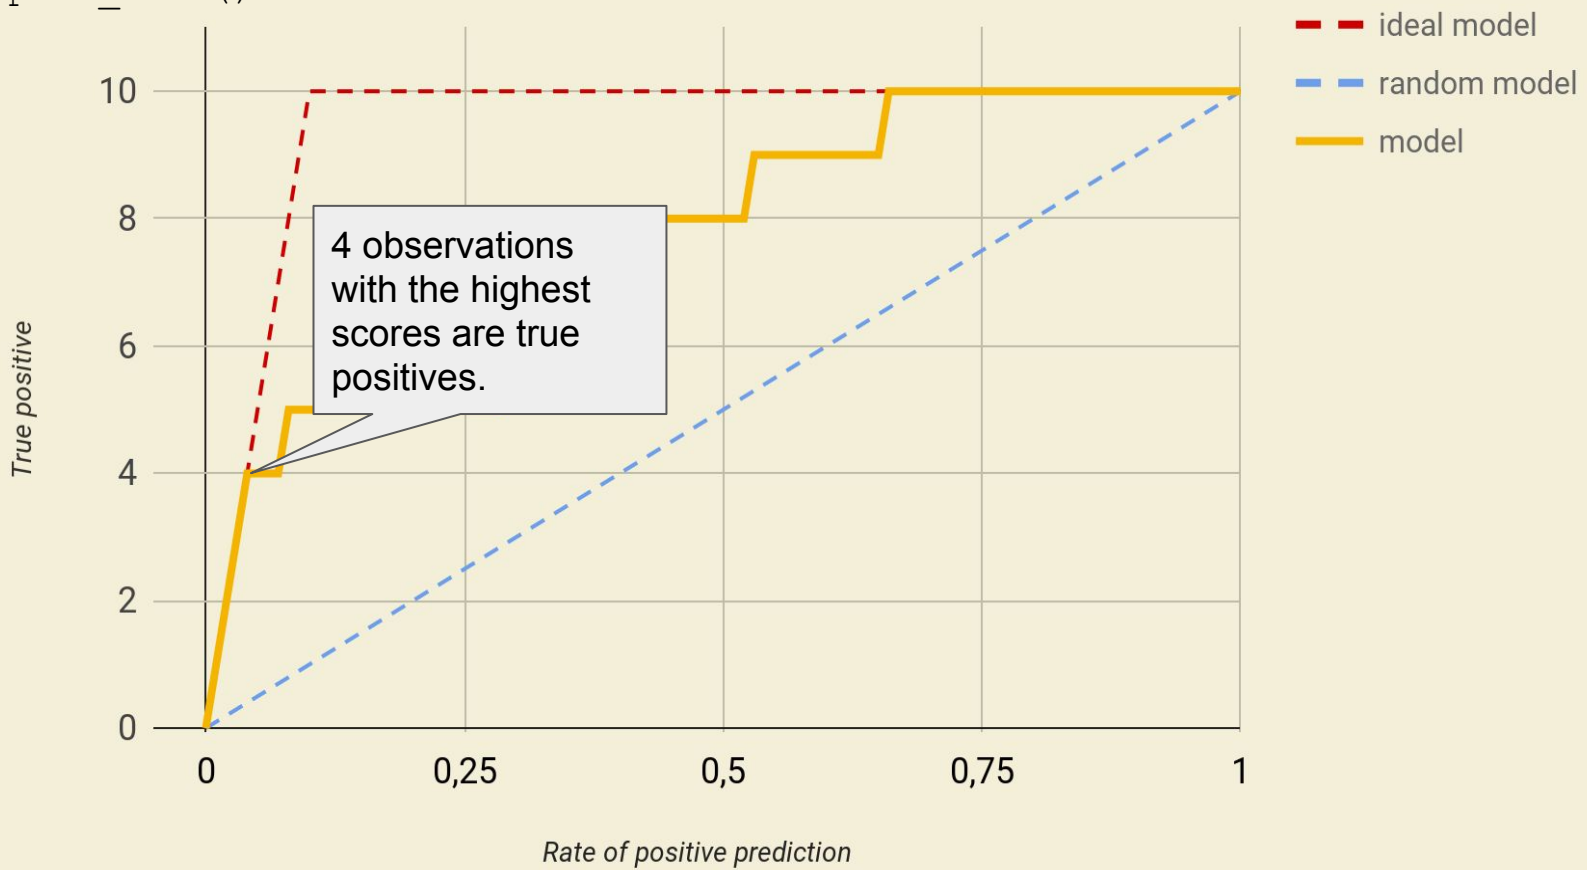

LIFT (toy example)

`plot_lift()`

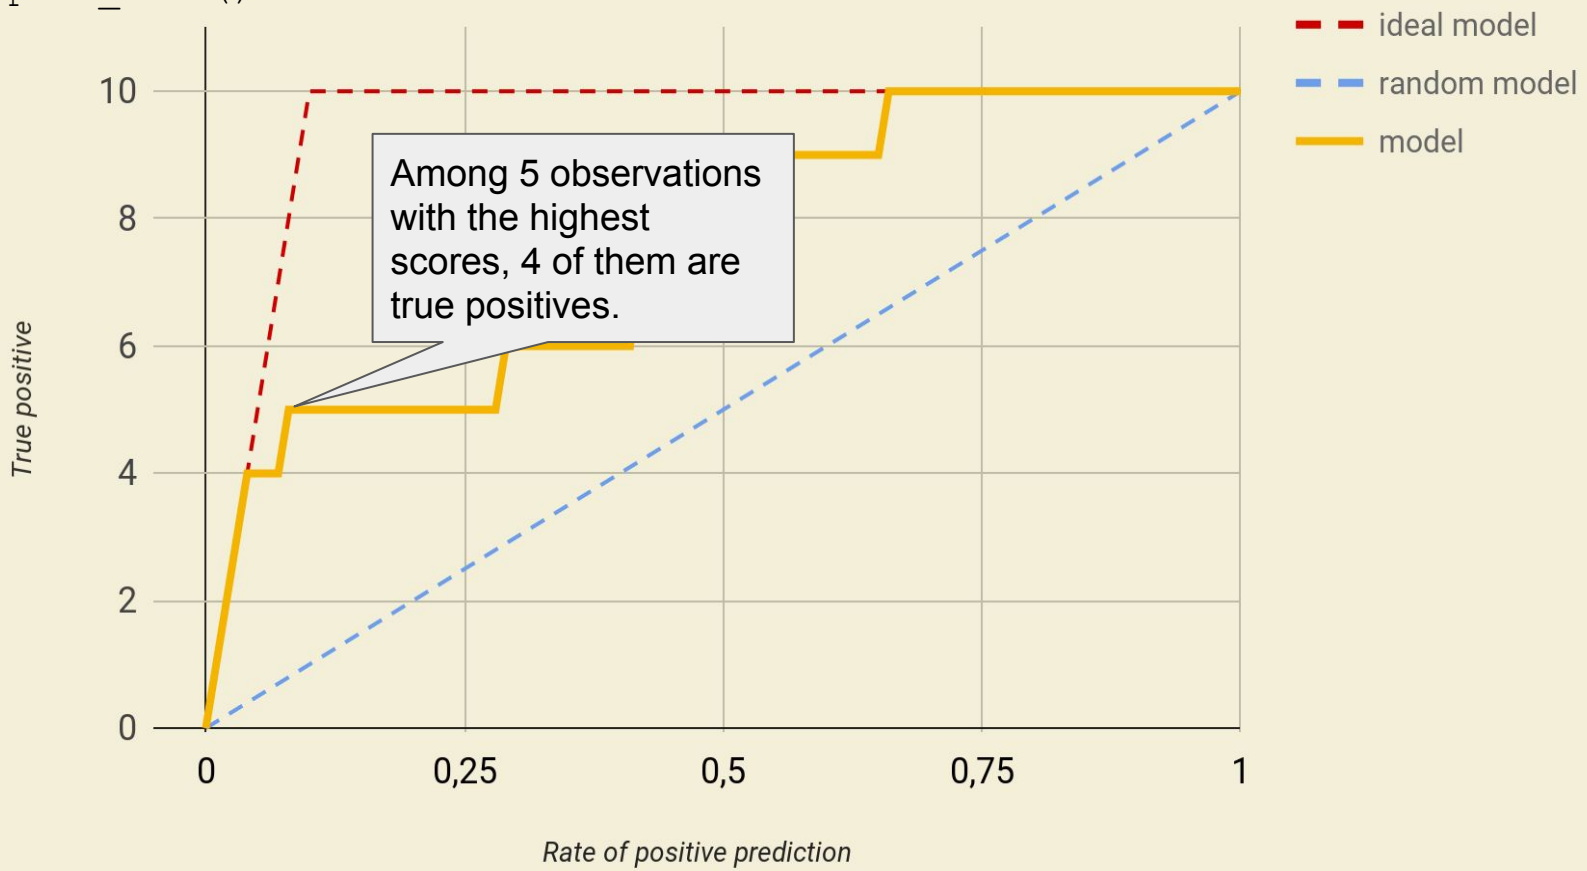

LIFT (toy example)

`plot_lift()`

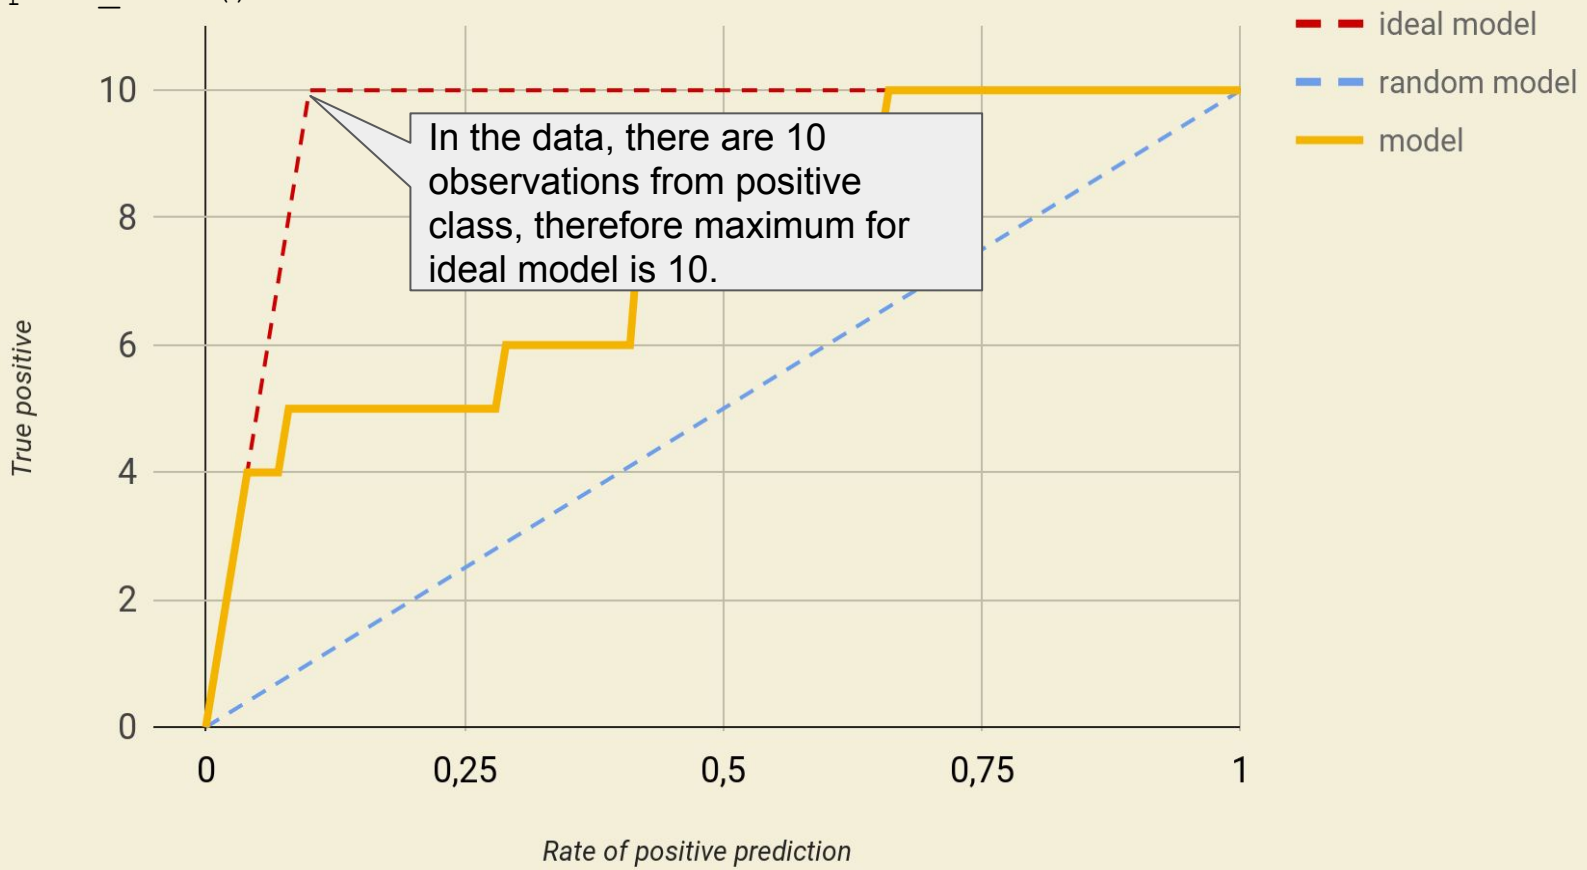

Feature importance

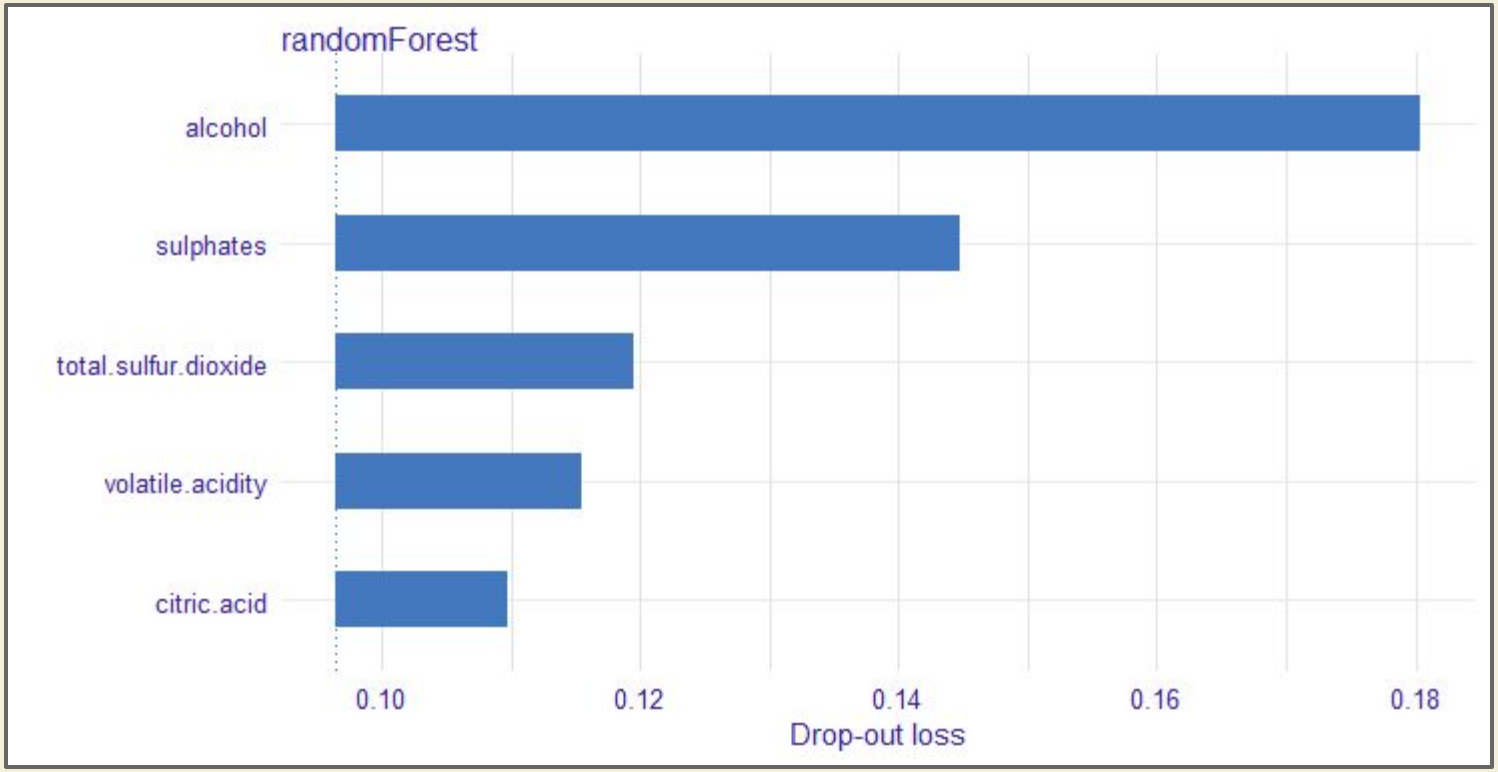

https://pbiecek.github.io/PM_VEE/featureImportance.html

Feature importance

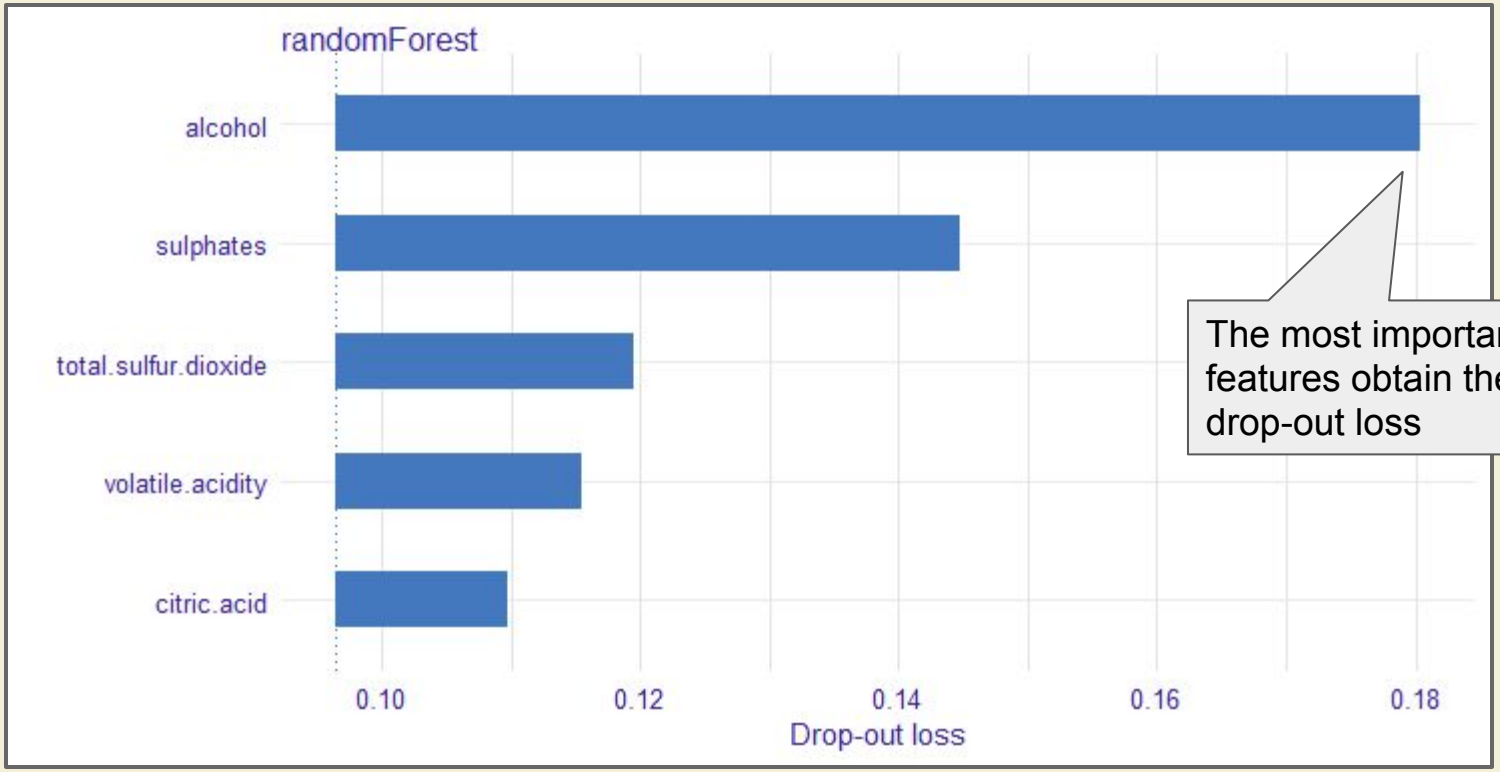

https://pbiecek.github.io/PM_VEE/featureImportance.html

Partial Dependence Plot

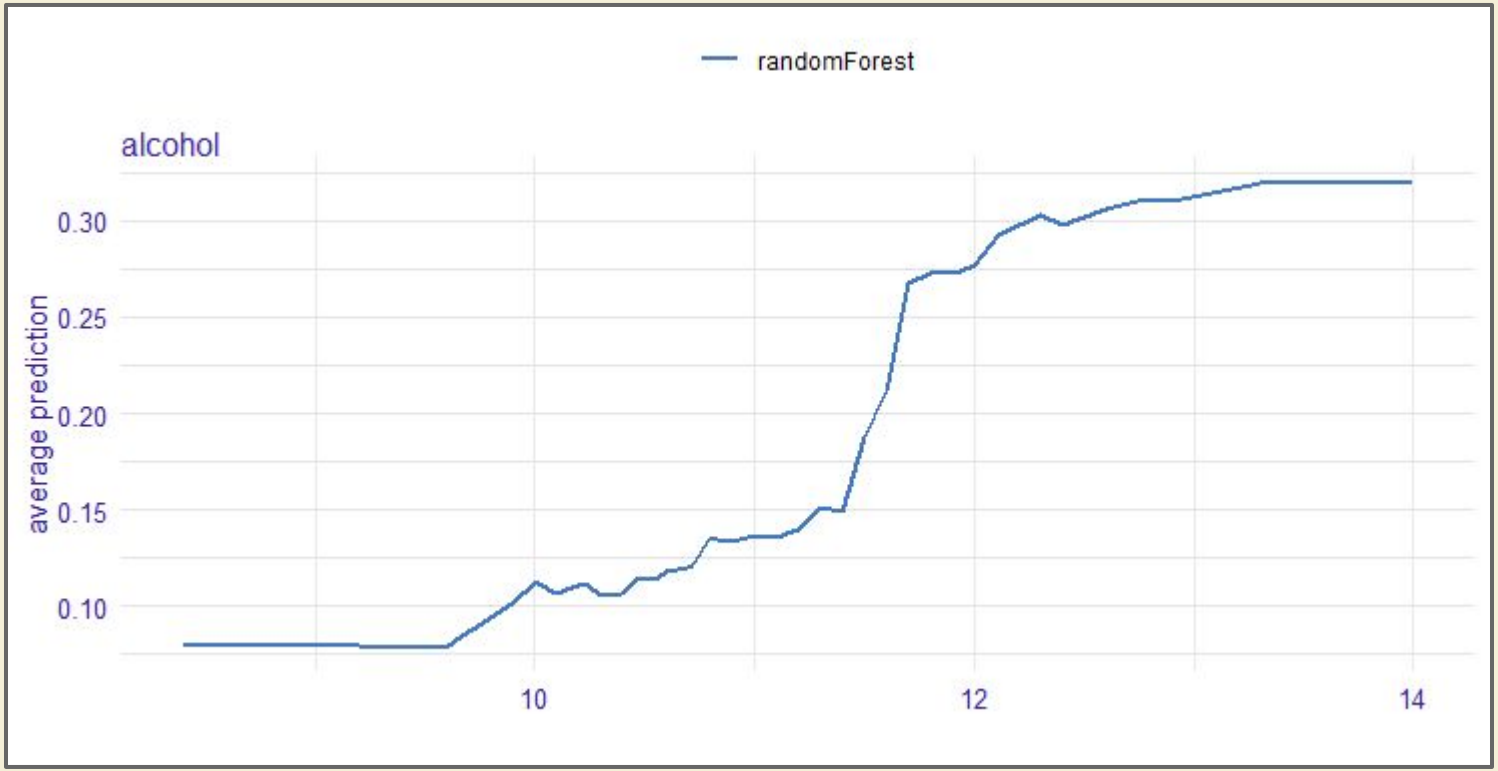

https://pbiecek.github.io/PM_VEE/partialDependenceProfiles.html

Partial Dependence Plot

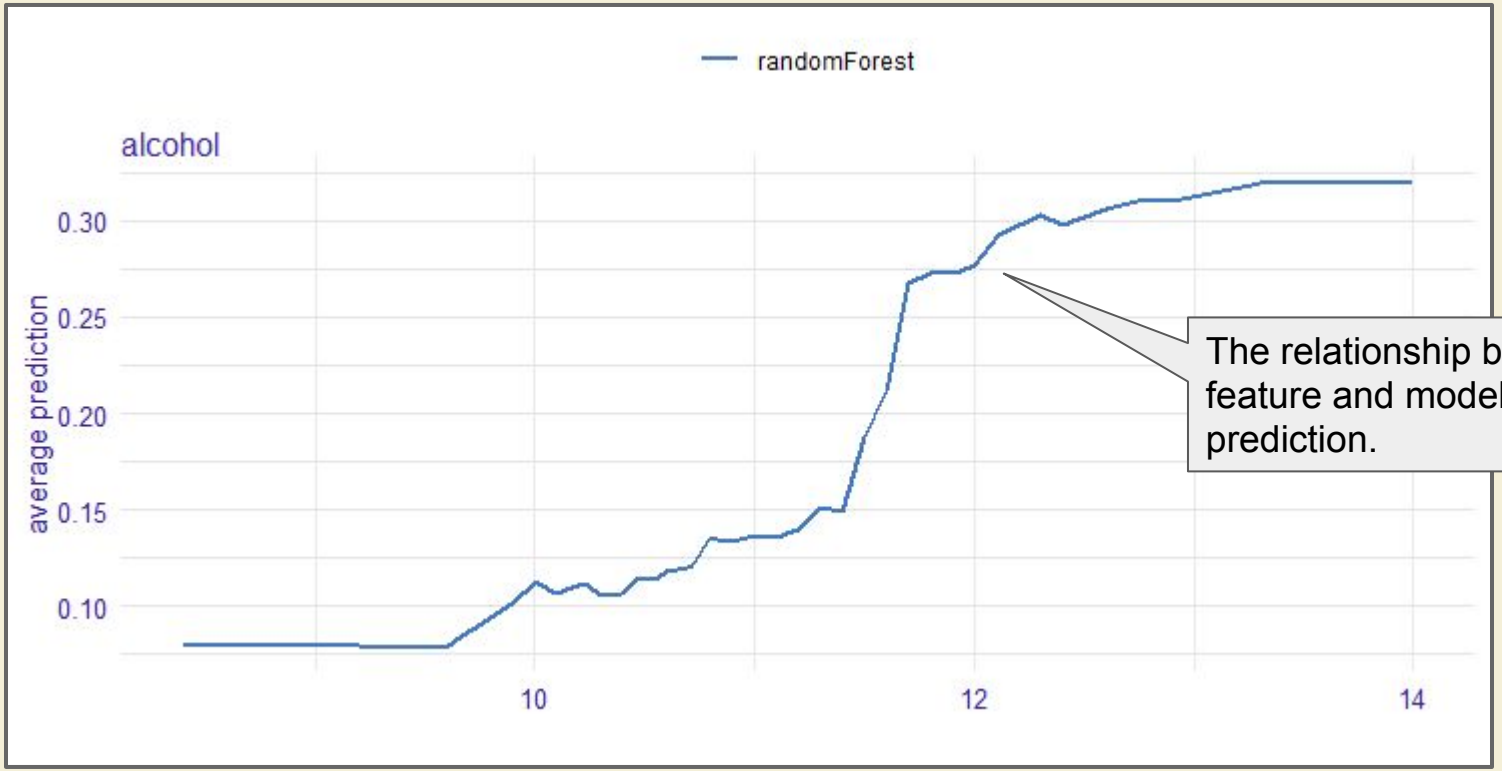

https://pbiecek.github.io/PM_VEE/partialDependenceProfiles.html

Ceteris Paribus Profile

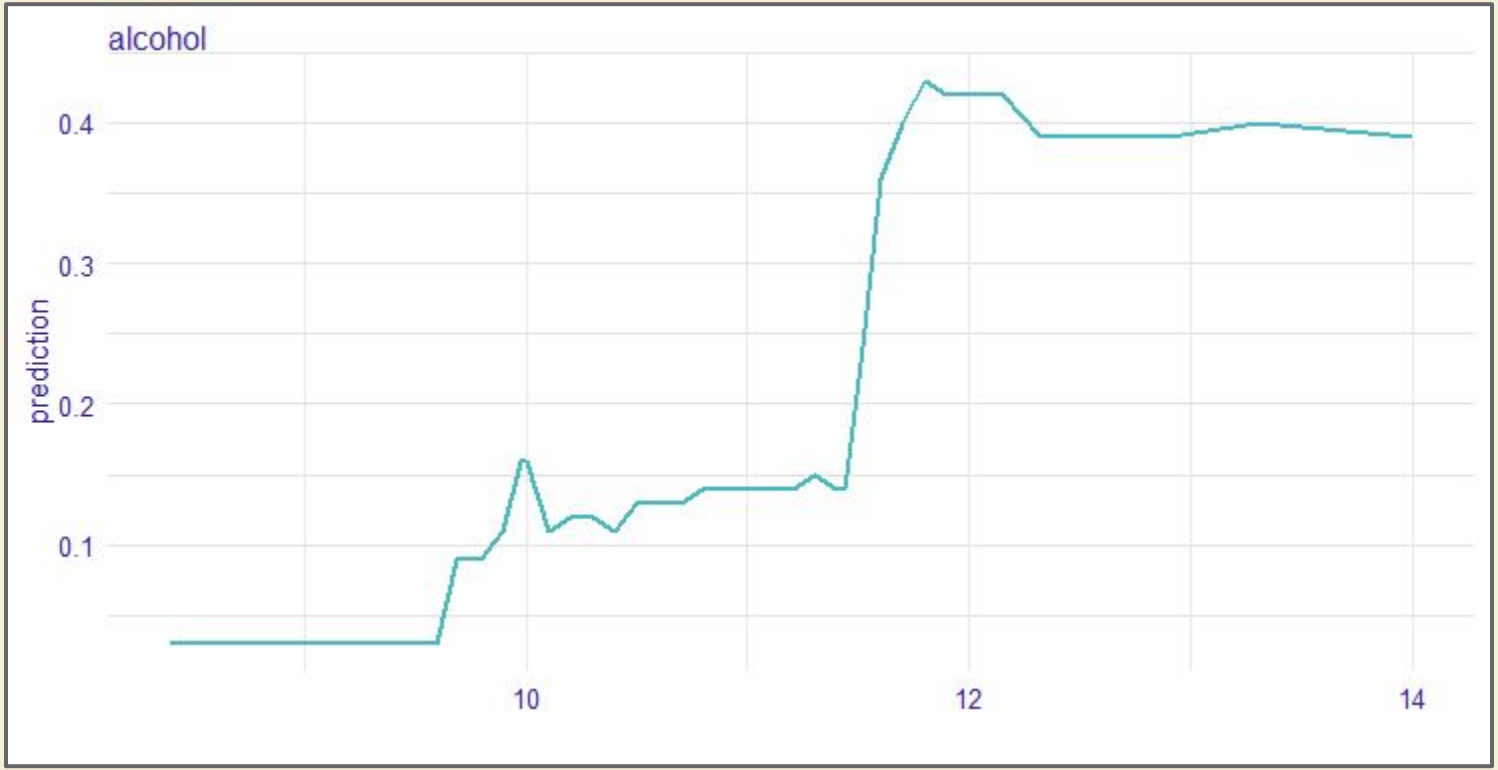

https://pbiecek.github.io/PM_VEE/ceterisParibus.html

Ceteris Paribus Profile

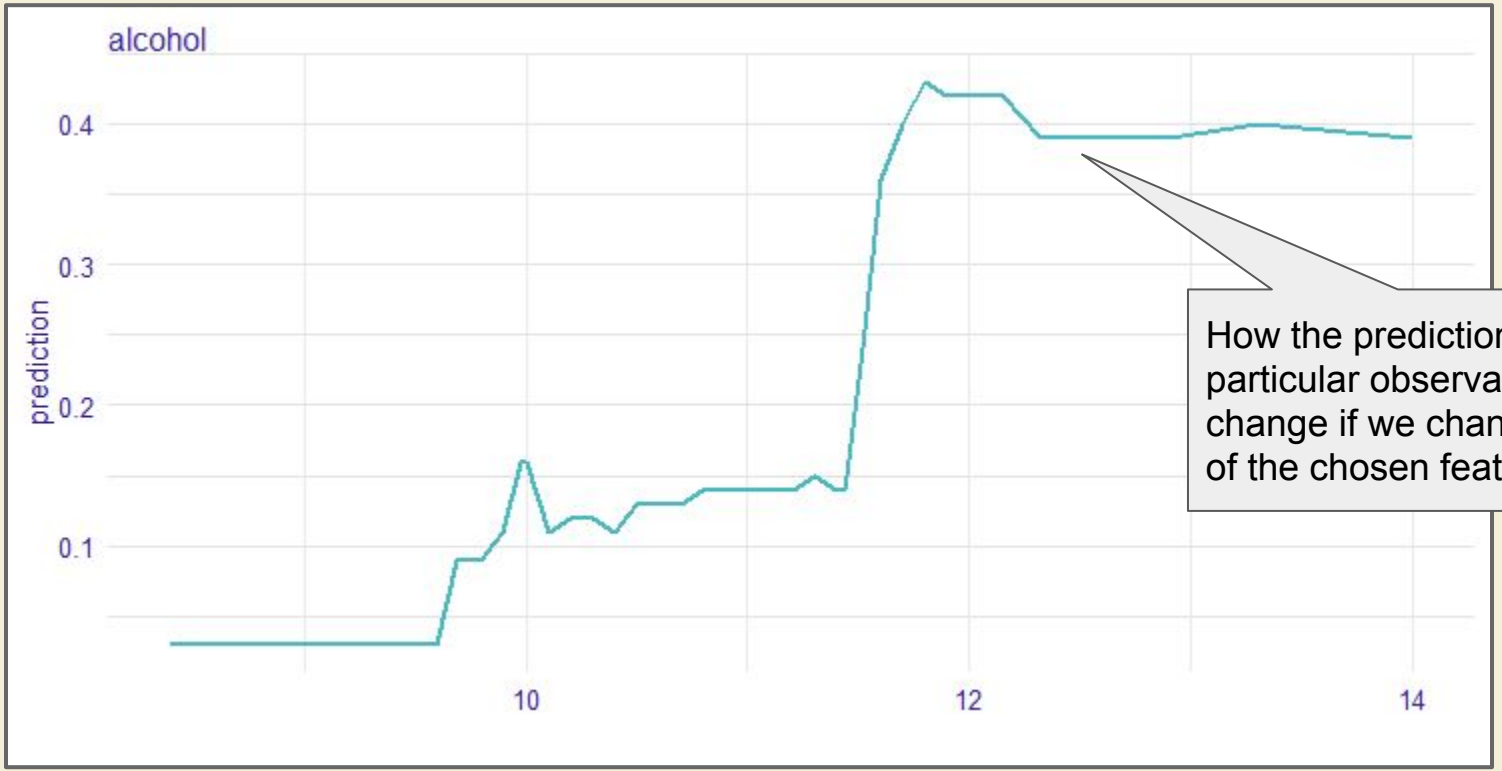

https://pbiecek.github.io/PM_VEE/ceterisParibus.html

Break Down

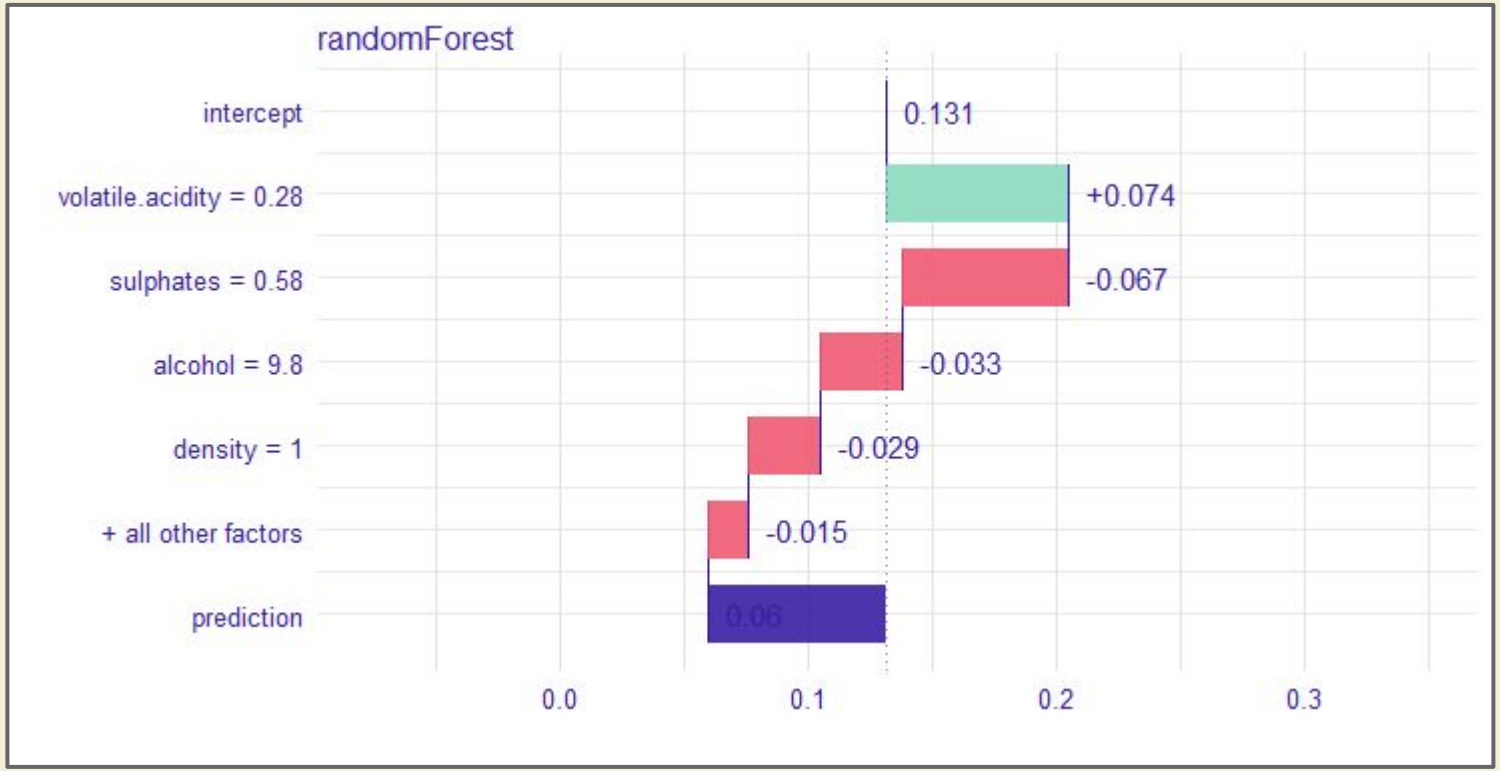

https://pbiecek.github.io/PM_VEE/breakDown.html

Break Down

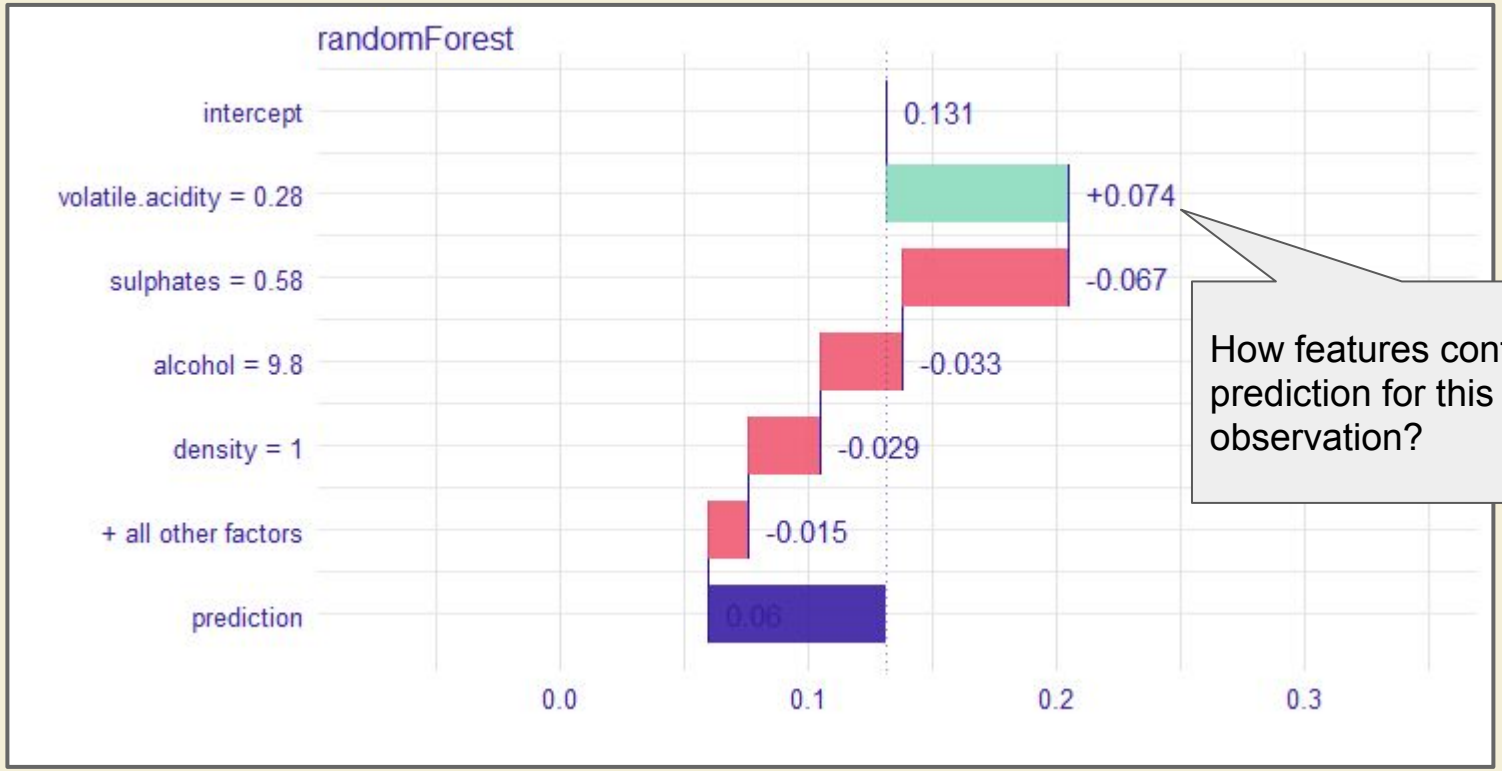

How features contribute into prediction for this particular observation?

https://pbiecek.github.io/PM_VEE/breakDown.html

Plot predictions (against target variable)

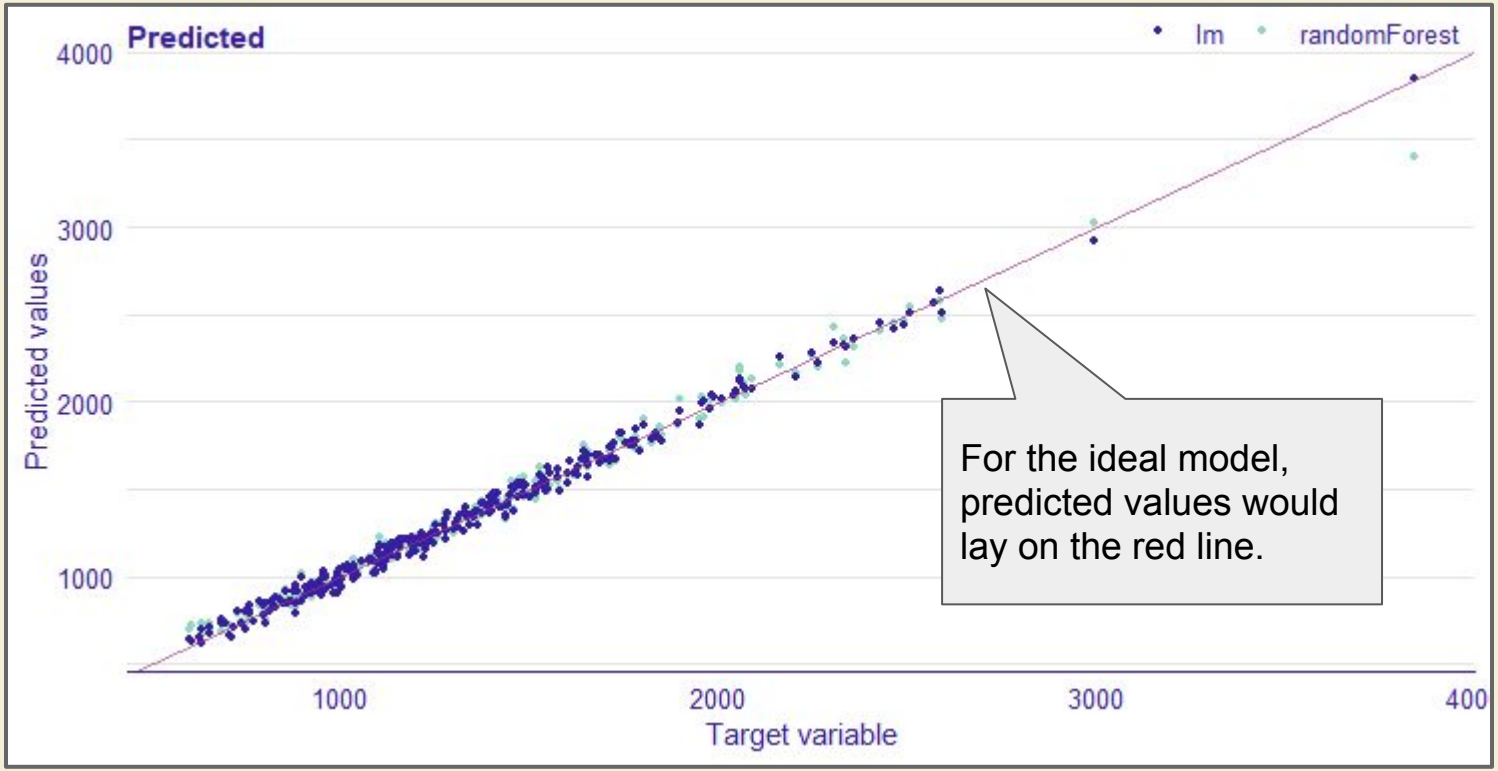

Plot predictions (against any variable)

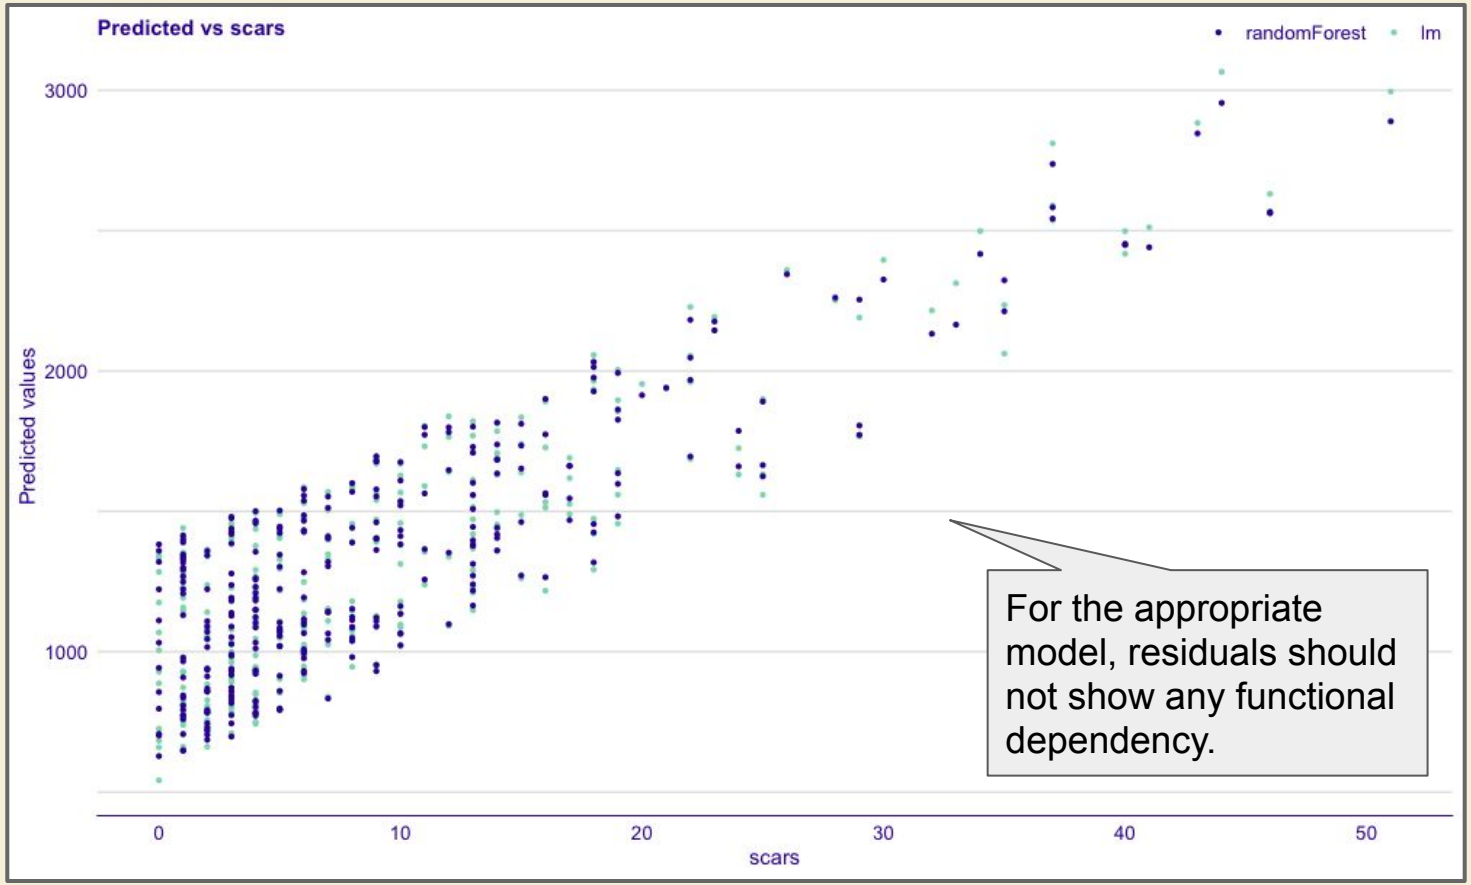

Residual Density

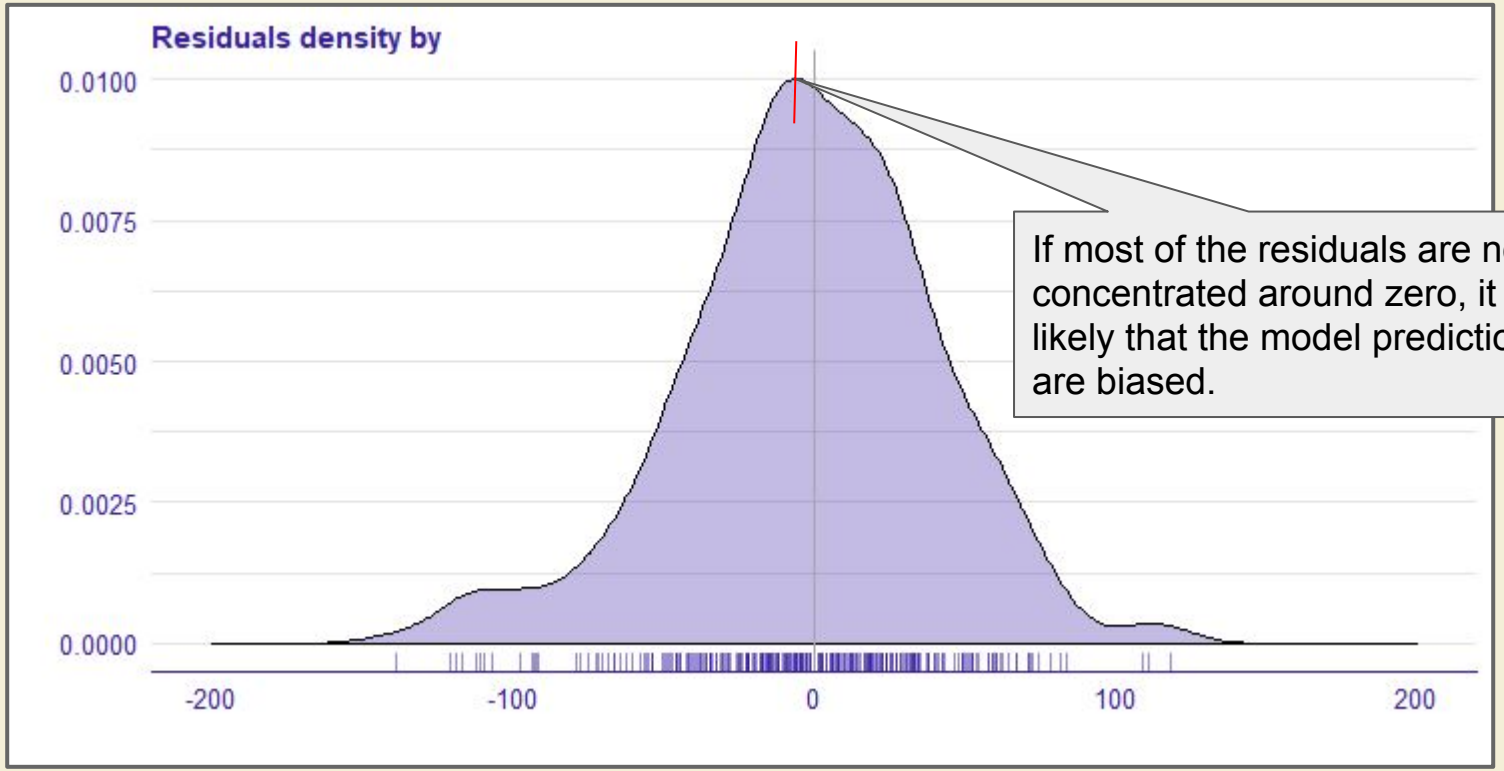

Data sets

Wine quality (1)

Data set used to predict human wine taste preferences (part of dataset of red and white variants of the Portuguese "*Vinho Verde*" wine).

Contains physicochemical (inputs) and sensory (the output) variables (no data about grape types, wine brand, wine selling price, etc.).

The datasets can be viewed as classification or regression tasks.

Wine quality (2)

The classes are ordered and not balanced (e.g. there are many more normal wines than excellent or poor ones). Outlier detection algorithms could be used to detect the few excellent or poor wines.

Interesting in testing feature selection methods, as it is not clear if all input variables are relevant.

The details are described in: P. Cortez, A. Cerdeira, F. Almeida, T. Matos and J. Reis. *Modeling wine preferences by data mining from physicochemical properties*. In Decision Support Systems, Elsevier, 47(4):547-553. ISSN: 0167-9236.

Wine quality - variable description (1)

Fixed acidity: most acids involved with wine or fixed or nonvolatile (do not evaporate readily)

Volatile acidity: the amount of acetic acid in wine, which at too high of levels can lead to an unpleasant, vinegar taste

Citric acid: found in small quantities, citric acid can add ‘freshness’ and flavor to wines

Residual sugar: the amount of sugar remaining after fermentation stops, it’s rare to find wines with less than 1 gram/liter and wines with greater than 45 grams/liter are considered sweet

Chlorides: the amount of salt in the wine

Free sulfur dioxide: the free form of SO2 exists in equilibrium between molecular SO2 (as a dissolved gas) and bisulfite ion; it prevents microbial growth and the oxidation of wine

Wine quality - variable description (2)

Total sulfur dioxide: amount of free and bound forms of S02; in low concentrations, SO2 is mostly undetectable in wine, but at free SO2 concentrations over 50 ppm, SO2 becomes evident in the nose and taste of wine

Density: the density of water is close to that of water depending on the percent alcohol and sugar content

pH: describes how acidic or basic a wine is on a scale from 0 (very acidic) to 14 (very basic); most wines are between 3-4 on the pH scale

Sulphates: a wine additive which can contribute to sulfur dioxide gas (S02) levels, which acts as an antimicrobial and antioxidant

Alcohol: the percent alcohol content of the wine

Quality: output variable (based on sensory data, score between 0 and 10)

Dragons data

Values are generated in a way to: have nonlinearity in **year_of_birth** and height and have **concept drift** in the test set

year_of_birth:	year in which the dragon was born. Negative year means year BC, eg: -1200 = 1201 BC
year_of_discovery:	year in which the dragon was found.
height:	height of the dragon in yards.
weight:	weight of the dragon in tons.
scars:	number of scars.
colour:	colour of the dragon.
number_of_lost_teeth:	number of teeth that the dragon lost.
life_length:	life length of the dragon.

MI2DataLab research lab is looking for you!

Are you intreseted in XAI, AutoML, AutoEDA other innovations in the next generation of ML?

Come to us
http://bit.do/MI2isHiring

Appendix

Confusion matrix

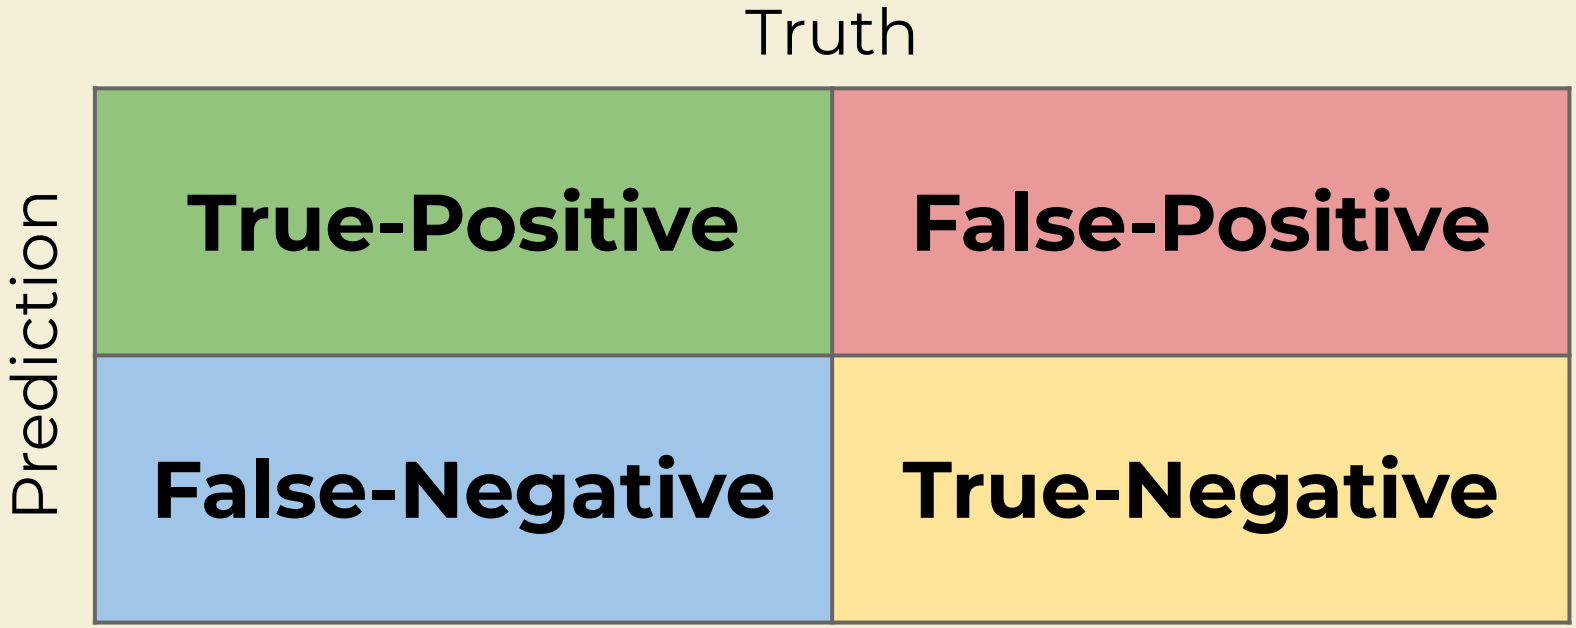

Confusion matrix - some scores

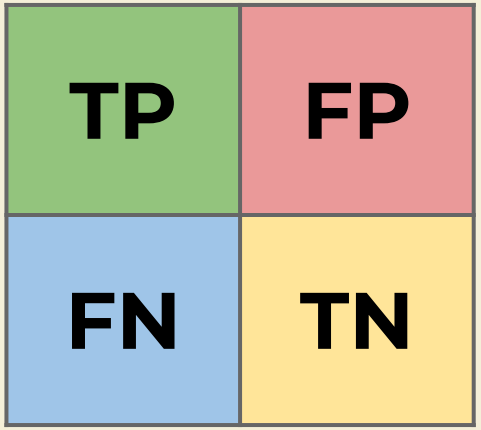

Accuracy: (TP + TN) / (TP + TN + FP + FN)

Recall: TP / (TP + FN)

Precision: TP / (TP + FP)

F-measure: (2 * Recall * Precision) / (Recall + Precision)